

Results

Synthetic paired pilot, batches A-E. Sample size is 5 paired observations. This is a document-handling acceptance fixture, not a product benchmark or certification.

Median baseline 205 ms	Median candidate 100 ms
Median paired saving 105 ms	Median paired reduction 50.0%

Baseline/candidate median ratio: **2.05x** (205 / 100).

Batch	Baseline (ms)	Candidate (ms)	Saving (ms)	Reduction
A	200	100	100	50.00%
B	210	105	105	50.00%
C	190	95	95	50.00%
D	205	100	105	51.22%
E	220	110	110	50.00%

All five candidate values are lower than their paired baseline values in this synthetic sample. This is descriptive only and is not a population success-rate estimate.

Methods and limitations

Method

All rows marked `include=true` in `observations.csv` were included. Medians are the middle values after sorting each five-value series.

Paired saving is baseline minus candidate. Paired percentage reduction is $100 \times (\text{baseline} - \text{candidate}) / \text{baseline}$. The reported 50.0% is the median of the five paired percentage reductions. The workbook ratio is the ratio of the two series medians: $205 / 100 = 2.05$.

Workbook correction

`metrics.xlsx` was inspected before editing. Raw elapsed-time measurements were preserved. Batch E's inclusion flag was changed from false to true to match the current source CSV. The Summary was updated to:

- `sample_n = 5`
- `median_baseline_ms = 205`
- `median_candidate_ms = 100`
- `ratio = 2.05`

Independent verification

A separate local Python calculation from `observations.csv` returned `n=5`, baseline median 205 ms, candidate median 100 ms, paired-difference median 105 ms, and paired-reduction median 50.0%.

Limitations

- Only five paired observations are present.
- The data are synthetic and intended for document acceptance testing.
- No confidence interval, population success rate, generalization claim, production-performance claim, or certification is warranted.